

How to use this template

1. Copy this file into the dataset folder or repository root and fill the YAML front-matter.
2. Keep the ``manifest_file`` and ``checksum`` up to date when files change.
3. Link the ``commit`` field to the code that produces the dataset so others can reproduce it.

Minimal example

```
---
dataset_name: "brain_mets_clinical"
dataset_id: "brain_mets_clinical_v1"
dataset_version: "1.0"
description: |
  Clinical and genomic data for brain metastasis
  patients collected from internal registry.
source: "Internal registry"
source_url: ""
date_acquired: "2024-12-15"
license: "internal-research-use"
access_restrictions: "controlled"
sensitive: true
data_root: "data/raw"
files:
  - path: "data/raw/clinical.csv"
    filename: "clinical.csv"
    format: "csv"
    size_bytes: 1234567
    checksum: "sha256:abcdef..."
    rows: 1024
schema: |
  - name: patient_id
    type: string
    description: "de-identified patient identifier"
    nullable: false
  - name: age
    type: integer
    description: "age at diagnosis"
    nullable: false
key_identifiers: ["patient_id"]
primary_key: "patient_id"
sample_size:
```

```
observations: 1024
units: "patients"
date_range:
  start: "2010-01-01"
  end: "2023-12-31"
update_frequency: "as-needed"
preprocessing_notes: |
  - Drop duplicate patient rows keeping first
  - Encode categorical variables with label encoding
missing_value_strategy: "impute_median_for_numeric,
mode_for_categorical"
provenance: |
  - created_by: "src/data collection and processing/
  ingest_clinical.py"
  - created_on: "2024-12-15"
  - commit: "abc123def"
contact:
  name: "Data Owner"
  email: "owner@example.org"
  role: "data steward"
manifest_file: "manifests/brain_mets_manifest.csv"
notes: |
  - Sensitive clinical variables present; dataset access
  controlled.
```

- See manuscript for cohort definition.

```
tags: ["clinical", "genomics", "brain-metastasis"]  
---
```


